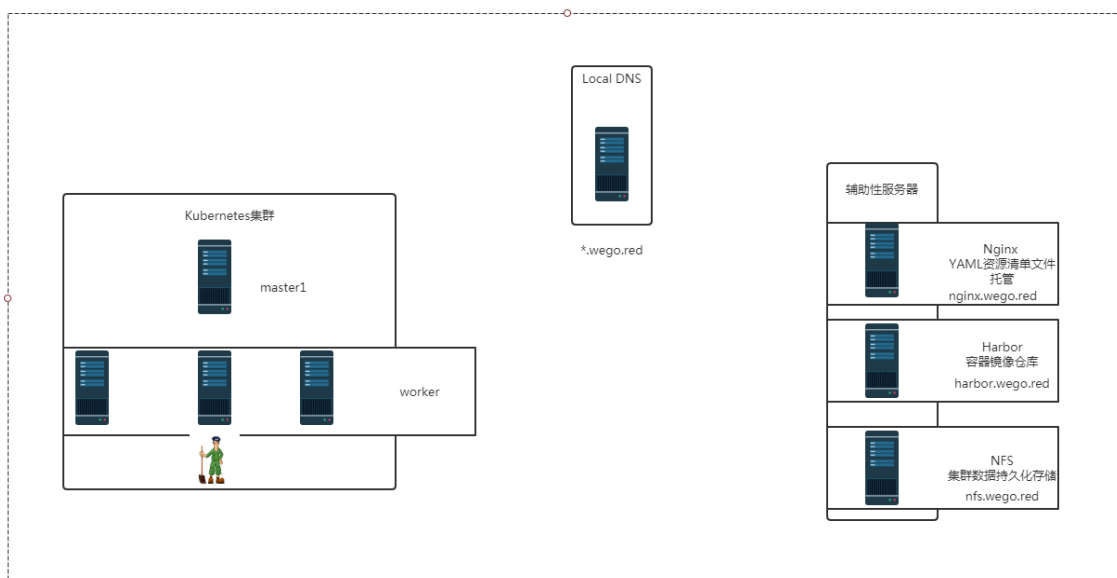
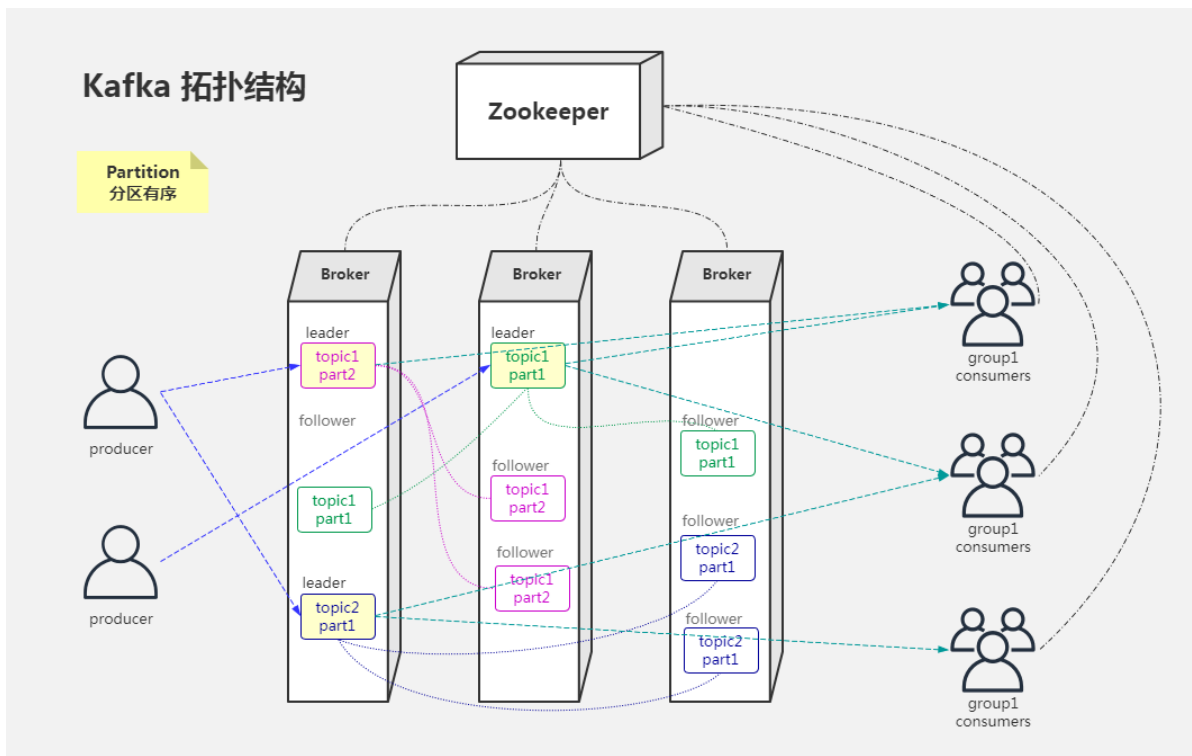


kubernetes集群 应用实践 kafka部署

零.1、环境说明



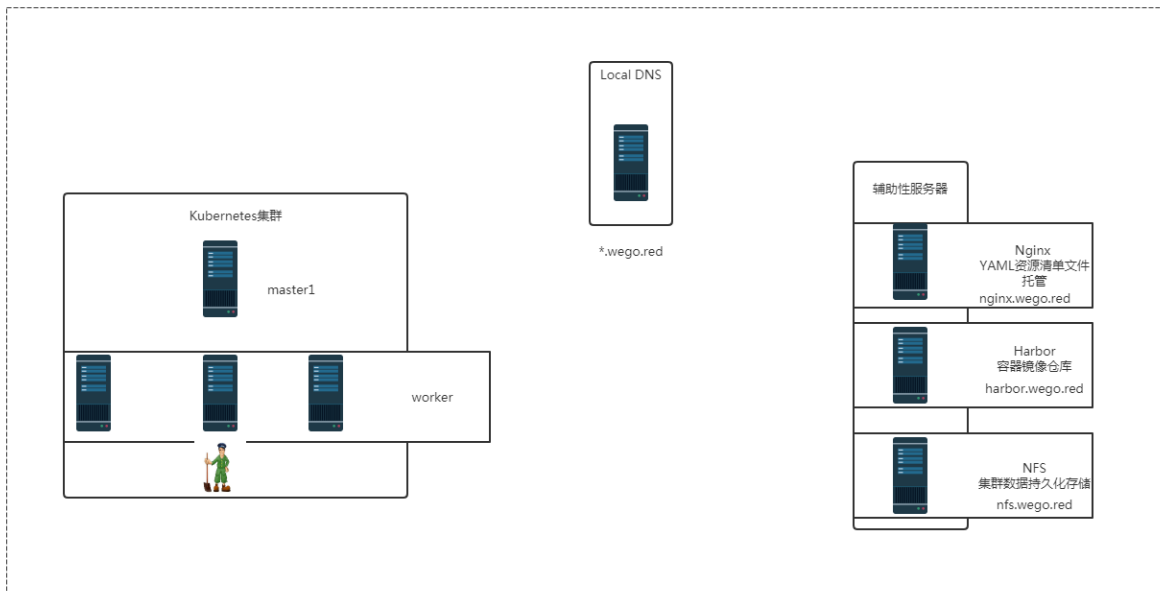
零.2、kafka架构说明



zookeeper在kafka集群中的作用

- 一、Broker注册
- 二、Topic注册
- 三、Topic Partition选主
- 四、生产者负载均衡
- 五、消费者负载均衡

一、持久化存储资源准备



1.1 创建共享目录

```
[root@nfsserver ~]# mkdir -p /vdc/kafka/data1
[root@nfsserver ~]# mkdir -p /vdc/kafka/data2
[root@nfsserver ~]# mkdir -p /vdc/kafka/data3
```

1.2 验证共享目录

```
[root@nfsserver ~]# tree /vdc/kafka
/vdc/kafka
├── data1
├── data2
└── data3
```

3 directories, 0 files

1.3 共享存储目录

```
[root@nfsserver ~]# cat /etc/exports
/vdc/zk/data1 *(rw, sync, no_root_squash)
/vdc/zk/data2 *(rw, sync, no_root_squash)
/vdc/zk/data3 *(rw, sync, no_root_squash)
/vdc/kafka/data1 *(rw, sync, no_root_squash)
/vdc/kafka/data2 *(rw, sync, no_root_squash)
/vdc/kafka/data3 *(rw, sync, no_root_squash)
```

```
[root@nfsserver ~]# exportfs -a
```

```
[root@nfsserver ~]# showmount -e 192.168.122.250
Export list for 192.168.122.250:
/vdc/kafka/data3 *
/vdc/kafka/data2 *
/vdc/kafka/data1 *
/vdc/zk/data3 *
/vdc/zk/data2 *
/vdc/zk/data1 *
```

二、k8s集群中kafka持久化存储PV准备

2.1 创建PV资源清单文件

```
[root@nginxk8syaml kafka]# cat kafka-pv.yaml
apiVersion: v1
kind: PersistentVolume
metadata:
  name: kafka-data1
spec:
  capacity:
    storage: 3Gi
  accessModes:
    - ReadWriteOnce
  nfs:
    server: nfs.wego.red
    path: /vdc/kafka/data1
---
apiVersion: v1
kind: PersistentVolume
metadata:
  name: kafka-data2
spec:
  capacity:
    storage: 3Gi
  accessModes:
    - ReadWriteOnce
  nfs:
    server: nfs.wego.red
    path: /vdc/kafka/data2
---
apiVersion: v1
kind: PersistentVolume
metadata:
  name: kafka-data3
spec:
  capacity:
    storage: 3Gi
  accessModes:
    - ReadWriteOnce
```

```
nfs:
  server: nfs.wego.red
  path: /vdc/kafka/data3
```

2.2 定位资源清单文件位置



2.3 应用资源清单文件

```
[root@master1 ~]# kubectl apply -f
http://nginxk8syaml.k8sonline.com.cn/kafka/kafka-pv.yaml
persistentvolume/kafka-data1 created
persistentvolume/kafka-data2 created
persistentvolume/kafka-data3 created
```

2.4 验证PV创建情况

```
[root@master1 ~]# kubectl get pv
```

NAME		CAPACITY	ACCESS MODES	RECLAIM
POLICY	STATUS	CLAIM		
STORAGECLASS	REASON	AGE		
kafka-data1	Available	3Gi	RWO	Retain
	19s			
kafka-data2	Available	3Gi	RWO	Retain
	19s			
kafka-data3	Available	3Gi	RWO	Retain
	19s			

三、k8s集群中部署kafka

3.1 kafka资源清单文件

```
[root@nginxk8syam1 kafka]# cat kafka.yaml
```

```

---
apiVersion: v1
kind: Service
metadata:
  name: kafka-hs
  namespace: default
  labels:
    app: kafka
spec:
  ports:
    - port: 9093
      name: server
  clusterIP: None
  selector:
    app: kafka
---
apiVersion: policy/v1beta1
kind: PodDisruptionBudget
metadata:
  name: kafka-pdb
  namespace: default
spec:
  selector:
    matchLabels:
      app: kafka
  maxUnavailable: 1
---
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: kafka
  namespace: default
spec:

```

```

serviceName: kafka-hs
replicas: 3
selector:
  matchLabels:
    app: kafka
podManagementPolicy: Parallel
updateStrategy:
  type: RollingUpdate
template:
  metadata:
    labels:
      app: kafka
  spec:
    terminationGracePeriodSeconds: 300
    containers:
      - name: k8skafka
        imagePullPolicy: IfNotPresent
        image: harbor.wego.red/library/kubernetes-kafka:v10.2.1
        resources:
          requests:
            memory: "256Mi"
            cpu: "0.1"
        ports:
          - containerPort: 9093
            name: server
        command:
          - sh
          - -c
          - "exec kafka-server-start.sh /opt/kafka/config/server.properties --
override broker.id=${HOSTNAME##*-} \
  --override listeners=PLAINTEXT://:9093 \
  --override zookeeper.connect=zk-cs.default.svc.cluster.local:2181 \
  --override log.dir=/var/lib/kafka \
  --override auto.create.topics.enable=true \
  --override auto.leader.rebalance.enable=true \
  --override background.threads=10 \
  --override compression.type=producer \
  --override delete.topic.enable=false \
  --override leader.imbalance.check.interval.seconds=300 \
  --override leader.imbalance.per.broker.percentage=10 \
  --override log.flush.interval.messages=9223372036854775807 \
  --override log.flush.offset.checkpoint.interval.ms=60000 \
  --override log.flush.scheduler.interval.ms=9223372036854775807 \
  --override log.retention.bytes=-1 \
  --override log.retention.hours=168 \
  --override log.roll.hours=168 \
  --override log.roll.jitter.hours=0 \
  --override log.segment.bytes=1073741824 \
  --override log.segment.delete.delay.ms=60000 \
  --override message.max.bytes=1000012 \
  --override min.insync.replicas=1 \
  --override num.io.threads=8 \
  --override num.network.threads=3 \
  --override num.recovery.threads.per.data.dir=1 \
  --override num.replica.fetchers=1 \
  --override offset.metadata.max.bytes=4096 \
  --override offsets.commit.required.acks=-1 \
  --override offsets.commit.timeout.ms=5000 \

```

```

--override offsets.load.buffer.size=5242880 \
--override offsets.retention.check.interval.ms=600000 \
--override offsets.retention.minutes=1440 \
--override offsets.topic.compression.codec=0 \
--override offsets.topic.num.partitions=50 \
--override offsets.topic.replication.factor=3 \
--override offsets.topic.segment.bytes=104857600 \
--override queued.max.requests=500 \
--override quota.consumer.default=9223372036854775807 \
--override quota.producer.default=9223372036854775807 \
--override replica.fetch.min.bytes=1 \
--override replica.fetch.wait.max.ms=500 \
--override replica.high.watermark.checkpoint.interval.ms=5000 \
--override replica.lag.time.max.ms=10000 \
--override replica.socket.receive.buffer.bytes=65536 \
--override replica.socket.timeout.ms=30000 \
--override request.timeout.ms=30000 \
--override socket.receive.buffer.bytes=102400 \
--override socket.request.max.bytes=104857600 \
--override socket.send.buffer.bytes=102400 \
--override unclean.leader.election.enable=true \
--override zookeeper.session.timeout.ms=6000 \
--override zookeeper.set.acl=false \
--override broker.id.generation.enable=true \
--override connections.max.idle.ms=600000 \
--override controlled.shutdown.enable=true \
--override controlled.shutdown.max.retries=3 \
--override controlled.shutdown.retry.backoff.ms=5000 \
--override controller.socket.timeout.ms=30000 \
--override default.replication.factor=1 \
--override fetch.purgatory.purge.interval.requests=1000 \
--override group.max.session.timeout.ms=300000 \
--override group.min.session.timeout.ms=6000 \
--override inter.broker.protocol.version=0.10.2-IV0 \
--override log.cleaner.backoff.ms=15000 \
--override log.cleaner.dedupe.buffer.size=134217728 \
--override log.cleaner.delete.retention.ms=86400000 \
--override log.cleaner.enable=true \
--override log.cleaner.io.buffer.load.factor=0.9 \
--override log.cleaner.io.buffer.size=524288 \
--override log.cleaner.io.max.bytes.per.second=1.7976931348623157E308 \

--override log.cleaner.min.cleanable.ratio=0.5 \
--override log.cleaner.min.compaction.lag.ms=0 \
--override log.cleaner.threads=1 \
--override log.cleanup.policy=delete \
--override log.index.interval.bytes=4096 \
--override log.index.size.max.bytes=10485760 \
--override log.message.timestamp.difference.max.ms=9223372036854775807 \

--override log.message.timestamp.type=CreateTime \
--override log.preallocate=false \
--override log.retention.check.interval.ms=300000 \
--override max.connections.per.ip=2147483647 \
--override num.partitions=3 \
--override producer.purgatory.purge.interval.requests=1000 \
--override replica.fetch.backoff.ms=1000 \
--override replica.fetch.max.bytes=1048576 \

```



```

--override replica.fetch.response.max.bytes=10485760 \
--override reserved.broker.max.id=1000 "
env:
- name: KAFKA_HEAP_OPTS
  value : "-Xmx256M -Xms256M"
- name: KAFKA_OPTS
  value: "-Dlogging.level=INFO"
volumeMounts:
- name: datadir
  mountPath: /var/lib/kafka
readinessProbe:
  exec:
    command:
      - sh
      - -c
      - "/opt/kafka/bin/kafka-broker-api-versions.sh --bootstrap-
server=localhost:9093"
volumeClaimTemplates:
- metadata:
  name: datadir
  spec:
    accessModes: [ "ReadWriteOnce" ]
    resources:
      requests:
        storage: 3Gi

```

3.1.1 修改容器镜像

在课程目录中有相关镜像，直接导入harbor主机并上传到harbor仓库

修改容器镜像

```

containers:
- name: k8skafka
  imagePullPolicy: IfNotPresent
  image: harbor.wego.red/library/kubernetes-kafka:v10.2.1

```

3.1.2 修改zookeeper连接地址

修改zookeeper.connect连接地址，使用k8s集群内域名，需要DNS服务。

```

command:
- sh
- -c
- "exec kafka-server-start.sh /opt/kafka/config/server.properties --
override broker.id=${HOSTNAME##*-} \
--override listeners=PLAINTEXT://:9093 \
--override zookeeper.connect=zkc.default.svc.cluster.local:2181 \

```

3.2 应用kafka资源清单文件

```
[root@master1 ~]# kubectl apply -f http://nginx.wego.red/kafka/kafka.yaml
service/kafka-hs created
poddisruptionbudget.policy/kafka-pdb created
statefulset.apps/kafka created
```

3.3 验证kafka创建情况

```
[root@master1 ~]# kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
busybox-pod	1/1	Running	397	28d
kafka-0	1/1	Running	0	106s
kafka-1	1/1	Running	0	106s
kafka-2	1/1	Running	0	106s
nfs-client-provisioner-5786f95795-54v4s	1/1	Running	4	9d
zok-0	1/1	Running	2	3h38m
zok-1	1/1	Running	1	3h38m
zok-2	1/1	Running	1	3h37m

四、通过zookeeper查看broker

```
[root@master1 ~]# kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
zok-0	1/1	Running	2	5h4m
zok-1	1/1	Running	1	5h3m
zok-2	1/1	Running	1	5h3m

```
[root@master1 ~]# kubectl exec -it zok-1 -n default bash
root@zok-1:/# zkCli.sh
Connecting to localhost:2181
welcome to ZooKeeper!

[zk: localhost:2181(CONNECTING) 0] ls /
[cluster, controller, controller_epoch, brokers, zookeeper, admin,
isr_change_notification, consumers, hello, config]


[zk: localhost:2181(CONNECTED) 1] ls /brokers
[ids, topics, seqid]
[zk: localhost:2181(CONNECTED) 2] ls /brokers/ids
[0, 1, 2]
[zk: localhost:2181(CONNECTED) 3] get /brokers/ids/0
```

```

{"listener_security_protocol_map":{"PLAINTEXT":"PLAINTEXT"},"endpoints":
["PLAINTEXT://kafka-0.kafka-
hs.default.svc.cluster.local:9093"],"jmx_port":-1,"host":"kafka-0.kafka-
hs.default.svc.cluster.local","timestamp":"1581591232561","port":9093,"version":
4}
cZxid = 0x200000001a
ctime = Thu Feb 13 10:53:52 UTC 2020
mZxid = 0x200000001a
mtime = Thu Feb 13 10:53:52 UTC 2020
pZxid = 0x200000001a
cversion = 0
dataVersion = 0
aclVersion = 0
ephemeralOwner = 0x1703e20e92c0000
dataLength = 254
numChildren = 0
[zk: localhost:2181(CONNECTED) 4] get /brokers/ids/1
{"listener_security_protocol_map":{"PLAINTEXT":"PLAINTEXT"},"endpoints":
["PLAINTEXT://kafka-1.kafka-
hs.default.svc.cluster.local:9093"],"jmx_port":-1,"host":"kafka-1.kafka-
hs.default.svc.cluster.local","timestamp":"1581591238213","port":9093,"version":
4}
cZxid = 0x200000001e
ctime = Thu Feb 13 10:53:58 UTC 2020
mZxid = 0x200000001e
mtime = Thu Feb 13 10:53:58 UTC 2020
pZxid = 0x200000001e
cversion = 0
dataVersion = 0
aclVersion = 0
ephemeralOwner = 0x1703e20e92c0001
dataLength = 254
numChildren = 0
[zk: localhost:2181(CONNECTED) 5] get /brokers/ids/2
{"listener_security_protocol_map":{"PLAINTEXT":"PLAINTEXT"},"endpoints":
["PLAINTEXT://kafka-2.kafka-
hs.default.svc.cluster.local:9093"],"jmx_port":-1,"host":"kafka-2.kafka-
hs.default.svc.cluster.local","timestamp":"1581591242262","port":9093,"version":
4}
cZxid = 0x2000000022
ctime = Thu Feb 13 10:54:02 UTC 2020
mZxid = 0x2000000022
mtime = Thu Feb 13 10:54:02 UTC 2020
pZxid = 0x2000000022
cversion = 0
dataVersion = 0
aclVersion = 0
ephemeralOwner = 0x2703e207a150000
dataLength = 254
numChildren = 0

```

五、kafka应用操作测试



5.1 查看pod

```
[root@master1 ~]# kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
busybox-pod	1/1	Running	398	28d
kafka-0	1/1	Running	0	87m
kafka-1	1/1	Running	0	87m
kafka-2	1/1	Running	0	87m

5.2 进入kafka pod

```
[root@master1 ~]# kubectl exec -it kafka-0 -n default bash
```

```
# pwd
```

```
/
```

```
# ls
```

```
KEYS  boot  etc   lib   media  opt   root  sbin  sys  usr
bin   dev   home  lib64  mnt    proc  run   srv   tmp  var
```

```
# cd /opt/kafka/bin
```

```
# pwd
```

```
/opt/kafka/bin
```

```
# ls
```

```
connect-distributed.sh      kafka-replica-verification.sh
connect-standalone.sh      kafka-run-class.sh
kafka-acls.sh               kafka-server-start.sh
kafka-broker-api-versions.sh kafka-server-stop.sh
kafka-configs.sh            kafka-simple-consumer-shell.sh
kafka-console-consumer.sh   kafka-streams-application-reset.sh
kafka-console-producer.sh   kafka-topics.sh
kafka-consumer-groups.sh   kafka-verifiable-consumer.sh
kafka-consumer-offset-checker.sh kafka-verifiable-producer.sh
kafka-consumer-perf-test.sh windows
kafka-mirror-maker.sh      zookeeper-security-migration.sh
kafka-preferred-replica-election.sh zookeeper-server-start.sh
kafka-producer-perf-test.sh zookeeper-server-stop.sh
kafka-reassign-partitions.sh zookeeper-shell.sh
kafka-replay-log-producer.sh
```

5.3 创建test topic

```
root@kafka-0:/opt/kafka/bin# ./kafka-topics.sh --create --topic test --zookeeper zk-cs.default.svc.cluster.local:2181 --partitions 3 --replication-factor 3
```

输出

Created topic "test".

5.4 查看 topic

```
root@kafka-0:/opt/kafka/bin# ./kafka-topics.sh --zookeeper zk-cs.default.svc.cluster.local:2181 --list
```

输出

test

5.5 模拟生产者

终端1

```
root@kafka-0:/opt/kafka/bin# ./kafka-console-producer.sh --topic test --broker-list kafka-0.kafka-hs.default.svc.cluster.local:9093,kafka-1.kafka-hs.default.svc.cluster.local:9093,kafka-2.kafka-hs.default.svc.cluster.local:9093
this is a test message
hello world
```

CTRL+C即可退出

5.6 模拟消费者

终端2

```
[root@master1 ~]# kubectl exec -it kafka-0 bash
root@kafka-0:/# cd /opt/kafka/bin
root@kafka-0:/opt/kafka/bin# ./kafka-console-consumer.sh --topic test --
zookeeper zk-cs.default.svc.cluster.local:2181 --from-beginning
Using the ConsoleConsumer with old consumer is deprecated and will be removed in
a future major release. Consider using the new consumer by passing [bootstrap-
server] instead of [zookeeper].
this is a test message
hello world
```

六、集群内业务系统访问kafka

```
broker-list:
kafka-0.kafka-hs.default.svc.cluster.local:9093
kafka-1.kafka-hs.default.svc.cluster.local:9093
kafka-2.kafka-hs.default.svc.cluster.local:9093
```